

Managing the Context Window in Code

Every request to Claude carries one context window, and it is a finite token budget you manage yourself.



ONE CONTEXT WINDOW

Finite token budget

0 TOKENS

IN USE

MAX TOKENS



	System Prompt	1,200
	User Message	8,400
	Code Snippet	15,600
	File Content	22,500
	Retrieved Context	31,800
	Conversation History	21,500

TOTAL IN USE

101,000 / 200,000



ONE CONTEXT WINDOW

Finite token budget

0 TOKENS

IN USE

MAX TOKENS




	System Prompt	1,200
	User Message	8,400
	Code Snippet	15,600
	File Content	22,500
	Retrieved Context	68,400 ⚠
	Conversation History	92,300 ⚠

OVER BUDGET

208,400 / 200,000

Three Foundations to Start



Context Window

The maximum tokens the model attends to in one request: system prompt, tool definitions, history, tool results, and the current generation.



Context Window Management

Deliberately controlling what occupies that finite budget so high-signal content stays in view. An active engineering task, not a model-managed behavior.



Attention Budget

Like working memory: every new token you add depletes it, so more context is not automatically more capability.

What Fills the Context Window

On a single request, more lands in the window than you'd guess.



The system prompt and every tool definition ride on each call

They consume tokens before your prompt even begins.



The full conversation history accumulates turn by turn

Every message you and the model exchange stays in the window.



Raw tool results are usually the single biggest source

APIs, files, and queries often return large payloads.

ONE CONTEXT WINDOW



System Prompt & Tool Definitions

You are a helpful assistant.
Tools: `get_weather(...)`, `search_web(...)`
Schemas, descriptions, constraints...



Conversation History

User: What's the forecast in Seattle?
Assistant: I'll check the weather.
Tool: `get_weather {"location": "Seattle"}`
Assistant: It's raining, 12C.
User: Compare with New York.
...



Tool Results (Raw Payloads)

```
{  
  "location": "Seattle, WA",  
  "temp_c": 12,  
  "condition": "Rain",  
  "humidity": 87,  
  "wind_kph": 18,  
  "forecast": [ ... 48 hourly items ... ]  
}
```



Model Output (Current Turn)

Assistant: In Seattle it's 12C and raining.
In New York it's 18C and partly cloudy...

Assumption vs. Reality



The Assumption

Many builders expect the model to quietly forget stale turns and keep only what matters, the way a person lets old details fade from memory over time.




The Reality

Accumulation is additive. Nothing is dropped unless your app prunes or compacts it. The API never auto-summarizes on overflow, so ownership sits with you.

Context Rot and the Attention Budget

Two forces explain why a stuffed window hurts.

Recall Accuracy vs. Token Count



Context rot: as tokens climb, accurate recall of any one fact drops



A buried instruction competes with thousands of low-value tokens



The attention budget is spent faster as every token relates to every other

Three Ways Long Sessions Degrade



Drift

The model loses the thread as tangents and stale results dilute your instructions, and it reverts to earlier assumptions.



Bloat

The window fills with low-value tokens and verbose dumps, raising cost and latency and crowding out room to reason.



Overflow

You run past the token limit. A hard error you prevent, since the API will not auto-summarize to save you.


Context Drift Is a Composition Problem

When the model ignores a recent constraint, the instinct is to scold it.



Adding a louder instruction just adds tokens to a diluted window

More words do not fix attention competition.



Drift is a composition problem, not a model defect

The window is crowded with outdated and irrelevant context.



The fix is to re-ground the goal and prune the stale noise

Refresh the objective and remove irrelevant context to restore signal.

NOISY CONTEXT TOKENS



NEW INSTRUCTION
Do X, using format Y.



**DRIFTED
OUTPUT**

Three More Terms to Know




High-Signal Context

Keeping only tokens that earn their place, the smallest set that still drives the outcome. Informative, yet tight.



Re-Grounding

Restating the current goal and key constraints mid-session so a diluted instruction regains its weight, without piling on a harsher prompt.



Windowing

Keeping only a recent slice of history and letting older turns fall away, a bounded alternative to summarizing everything.

Informative, Yet Tight

Curate the window on purpose, every turn, rather than filling and forgetting.



Focused Context Window

Just what's needed.



User

What's the weather in Seattle?



Assistant

Checking the current weather for Seattle now.



Tool Result

```
{ "location": "Seattle, WA",  
  "temp_c": 14, "condition": "Cloudy" }
```



Assistant

It's 14 °C and cloudy in Seattle.

≈ 4 messages



Bloated Context Window

Too much noise.



User

What's the weather in Seattle?



Assistant

Sure, I can help with that. (3 msgs)



User

By the way, can you also recommend a cafe? (6 msgs)



Assistant

Here are some cafe ideas... (8 msgs)



User

Actually, what's the weather? (4 msgs)



Tool Result

```
{ "location": "Seattle, WA",  
  "temp_c": 14, "condition": "Cloudy" } (2 msgs)
```



Assistant

It's 14 °C and cloudy in Seattle. (3 msgs)

≈ 26+ messages



Before appending a big result, keep only what the model needs downstream

Summarize, filter, or extract to preserve signal and stay on track.



When the thread slips, re-ground instead of nagging

Re-center with a fresh query or context refresh to realign quickly.



Decide up front how history is bounded

Set clear rules for what to keep, summarize, or discard over time.

Owning the Context Window



It spans system prompt, tools, history, tool results, and generation



Nothing is dropped automatically; accumulation is additive




Drift and bloat come from an unmanaged, ever-growing window



The fix for drift is re-grounding and pruning, not a louder prompt

Pruning Tool Output and Compaction

In a tool-heavy agent loop, raw tool output is the fastest thing to fill your context window.




Three Ways to Reclaim Tokens




Tool-Output Pruning

Trimming, summarizing, or dropping verbose tool results so only the fields the model needs downstream survive. Your primary lever against bloat.



Context Compaction

Periodically compressing earlier turns into a shorter summary that keeps the essential thread while reclaiming tokens. A long session continues from the summary.



Tool-Result Clearing

A mechanical edit that swaps old tool results for placeholders while leaving your tool calls and reasoning intact. No summarization step.

Tool-Output Pruning Comes First

In a tool-heavy loop, pruning is usually the biggest single win.




One file read or **API response** can dump tens of thousands of tokens



Most of it is noise you will never reference again




Cut each result to its load-bearing fields, then drop the rest




Raw dumps are the fastest source of bloat, so trim them first.

Clearing vs. Compaction



Tool-Result Clearing

A mechanical server-side edit. It swaps old tool results for placeholders and keeps tool calls and reasoning. No model runs, so no inference cost, though the edit invalidates the prompt cache from that point.




Context Compaction

Runs a model to summarize older context, then continues from the summary. It costs an inference pass but can shrink dialogue and reasoning, not just results.

How Tool-Result Clearing Works

Clearing is configured as a context-management edit that watches your input tokens.

CONVERSATION TIMELINE



It fires once input crosses a set trigger threshold

When your context window usage exceeds the configured limit, clearing activates automatically.



It keeps the three most recent tool use and result pairs intact

The latest interactions stay complete so the model retains full conversational continuity.



Older results become placeholders; the model still sees the call ran

Placeholders preserve the record of what happened without consuming context space.


You can exclude specific tools so they are never cleared.

What a Compaction Pass Does



Summarizes

Compresses older turns into a concise summary and continues the session from it as the window fills.




Keeps the Gist

High-level facts, architectural decisions, and open questions central to the task tend to survive.



Drops Detail

Obscure specifics and exact verbatim phrasing are lost, so save anything you need word for word.



Costs Inference

It runs a model to build the summary, so unlike clearing it is not free.

Compaction Defaults and Cost


Compaction watches input tokens and triggers at a configured threshold.



The default trigger is 150K input tokens, with a floor of 50K



Once it fires, the API drops all content before the summary block



Summarizing runs a model, so it is not free like clearing

INPUT TOKENS

THRESHOLD



DROP BEFORE SUMMARY



GENERATE SUMMARY



KEEP SUMMARY BLOCK

Reach for it when the whole history has grown too large.

Three Clearing Controls



Keep

How many recent tool use and result pairs clearing preserves, so your current work is never wiped.



Trigger

The input-token threshold that decides when clearing fires. Cross it and the edit runs to free room.




Exclude Tools

A list of tools whose results are never cleared, used to protect output you always want visible.

Fresh Window vs. Compaction

When a session gets huge, compaction is not your only option.

Fresh Empty Window



```
main.py
```

```
1 |
2 |
3 |
4 |
5 |
6 |
```

Ln 1, Col 1 | Spaces: 4 | UTF-8 | Python

Compacted (Summarized History)



```
session_summary.md
```

```
1 # Session Summary
2
3 ## Completed
4 - Implemented user auth
5 - Added caching layer
6 - Fixed rate limiting bug
7
8 ## Current Focus
9 Improve search relevance
10
11 ## Next Steps
    - Add embedding index
    - Tune ranking weights
```

Ln 1, Col 1 | Spaces: 4 | UTF-8 | Markdown



A fresh context window can beat summarizing the old one



A capable agent rediscovers state from the git log, a progress file, or tests



Compact when the running thread genuinely matters

Reclaiming Tokens Deliberately



Prune verbose tool output first; it is your biggest win



Clearing is mechanical with no inference cost, keeping tool calls intact





Compaction runs a model and is lossy, but shrinks everything



Sometimes a fresh window beats compacting the old one

Isolating Context with Subagents

A long agent run fills the parent context with verbose discovery, file reads, and search dumps that crowd out the work that matters.



Three Terms to Anchor



Context Isolation

Delegating a scoped subtask to a subagent that returns only its result, keeping the subtask's intermediate tokens out of the parent's window.



Distilled Summary

The condensed result a subagent hands back, typically a thousand to two thousand tokens, not its full working trace.



Blast Radius

How far a bad intermediate step can spread, which isolation bounds by trapping the damage inside one subagent.

Why Isolation Saves Budget

The costliest tokens in an agent run are usually discovery, not the answer.



Reading files and searching dumps
huge raw output
into context



A subagent absorbs
that exploration in
its own window



The parent only
ever sees the
condensed result



Token Budget Comparison

Example: Log analysis task



Discovery / Explorer Tokens



Answer / Result Tokens



Without Isolation

Discovery
(120K tokens)

Answer
(8K tokens)



128K
Total

Full context window used



With Isolation

Subagent Explorer
(120K tokens)

Condensed Result
(4K tokens)

Parent Agent

4K tokens

Tiny context used by parent



Isolation cuts token spend dramatically.

Only the final summary reaches the parent.

Full Trace vs. Summary



Inside the Subagent

It may read dozens of files and run many searches, quietly building a long working trace of intermediate reasoning and tool output.




What Returns to the Parent

Only the distilled summary crosses back. That returned summary is the contract you build on, not the full trace the subagent worked through.

A Subagent Starts Fresh

A subagent does not inherit the parent's conversation. It sees only the task you hand it.



Anything it needs
goes into the delegated
prompt explicitly



Assumed shared context
is the most common
handoff bug




Scope the task tightly
so it knows exactly
what to do

Three Signs to Delegate



Verbose Discovery

The subtask reads or searches a lot and dumps output the parent never needs to keep.



Independent Workstream

The step can run on its own without sharing state back and forth with the main agent.



Bounded Blast Radius

You want a shaky intermediate step contained so its mistakes never reach the parent context.

The Summary Is the Contract

Treat what the subagent returns as its contract, not a window into how it got there.



Validate the summary
at the parent before
you act on it



Do not reach back
into the subagent's
working trace



Collect results
and synthesize them
into the next step




Parent Agent

Sets goal and
requests work

Interface

Subagent

Does work and
produces summary



Contract: Summary Only

```
{  
  "task_id": "analyze_sales_q1",  
  "status": "completed",  
  "findings": [  
    "Q1 revenue up 12% YoY",  
    "Top driver: product B"  
  ],  
  "metrics": {  
    "revenue_usd": 1245000,  
    "growth_pct": 12.3  
  },  
  "confidence": 0.87  
}
```



Validate



Aggregate



Synthesize



Next Step

Three Anti-Patterns to Avoid




Overly Broad Scope

Handing the subagent a fuzzy, sprawling task so it wanders and returns a summary you cannot trust.



Starved Context

Under-powering the subtask by withholding facts it needs, so it works blind and produces a weak result.




Parallel-Only Thinking

Treating subagents as mere concurrent workers while ignoring that isolation is their real value.

Isolate or Continue Inline

Not every subtask earns a subagent. Choose by how much noise it makes and whether its details matter later.



High intermediate noise
plus a throwaway trace
favors isolation



Details you reference
again later favor
staying inline



A quick single lookup
is faster done directly

Isolation Is the Point



A subagent isolates verbose work in its own window



Only a distilled summary returns to the parent





It starts fresh, so hand it what it needs




Treat that summary as the contract and validate it

Splitting Work Into Steps

One giant session that does everything accumulates noise until focus degrades.



Three Terms to Anchor



Workflow vs. Agent

A workflow is a developer-authored sequence with fixed control flow; an agent lets the model choose the next action at runtime.



Prompt Chaining

Breaking a task into sequential model calls, where each step's output feeds the next, so you can inspect and control the pipeline.



Staged Handoff

Passing forward only the distilled output of each stage, not its full working context, so the next step starts focused.

Session or Staged Steps

You can run a long job as one growing session or as a chain of focused steps.



A single session

carries every prior turn forward, noise and all



Staged steps

each start against a tight, relevant context



Each stage

passes on its distilled output, not its scratch work

One long session (bloated over time)



Staged steps (focused and passing outputs)



Same Call, Now Inside a Chain



Isolate It

The same call as last lesson: heavy throwaway noise and no shared state favor a subagent that returns only the answer.




Keep It Inline

Small steps whose details matter downstream, like a rename later steps touch, stay in the chain with state carried forward.

Pass Forward Distilled Output

The link that makes a workflow clean is what each step hands to the next.



Forward the distilled result, not the whole working trace

Pass only the compact facts needed for the next step to do its job.



Drop scaffolding the later steps will never reference

Remove debug logs, intermediate thoughts, and transient context.



Give each stage only the inputs its job requires


Provide the minimum fields and nothing the stage does not need.

Three Reasons to Chain Explicitly



Inspect Intermediate Output

You want to see and check a step's result before the next step runs on it.



Enforce Pipeline Structure

You need a specific, repeatable sequence of stages rather than letting the model improvise the path.



Enable Self-Correction

You run draft, then review, then refine as separate calls so each stage can catch the last one's errors.

The Self-Correction Loop

The most common chaining pattern is a model checking its own work in stages.




Draft: generate a first answer to the task



Review: check that draft against explicit criteria



Refine: fix the issues the review surfaced



Separate calls let you log and branch at each stage.

Three Handoff Habits




Distill Before Passing

Summarize each stage down to the load-bearing output before the next step ever sees it.



Validate Between Stages

Check a stage's output against your rules before the following step builds on it.



Keep Steps Independent

Design each stage so it needs only its inputs, not the entire history of the run.

Where Multi-Step Work Goes Wrong

A few habits quietly undo the benefit of splitting the work.



Anti-Pattern: Bloated Session Over-Chaining



Repeat 12+ times in one session



Problems

- ✗ Context grows huge and expensive
- ✗ Harder to debug and audit
- ✗ Small errors compound and are harder to recover



Slow, costly, brittle outcome



Letting one session grow instead of staging focused steps

Long sessions inflate context, increase cost, and make errors harder to isolate and fix.



Forwarding a full working trace instead of a distilled result

Passing every intermediate step creates noise, higher token usage, and distracts from what matters.



Over-chaining when the model's internal reasoning already handles it

Unnecessary hand-offs add latency and cost without improving outcomes.

Focused Steps Beat One Session



Run stages against tight context, not one growing session



Forward only the distilled output between stages




Chain explicitly to inspect output or enforce a pipeline



Use draft, review, refine to let steps self-correct

Writing Clear, Unambiguous Instructions

Clarity is the cheapest, highest-leverage reliability lever you have as a builder.



Three Marks of a Clear Ask



Specific

State exactly what you want. Name the output format, the fields, and the limits to respect.



Unambiguous

Leave no room for two readings. If a colleague could misread the ask, Claude can too.



Ordered

When order or completeness matters, give the steps as a numbered list Claude follows in sequence.

Treat Claude Like a New Hire

Think of Claude as a brilliant new employee who lacks your context.



Spell out norms and workflows you'd tell a new hire

Provide your processes, decision rules, escalation paths, and how work actually gets done.



Don't assume shared jargon, defaults, or house style

Define terms, acronyms, preferences, and formatting expectations up front.



Ask explicitly for above-and-beyond effort when you want it

State the standard, then specify when to go deeper, challenge assumptions, or add creativity.




Vague Ask vs. Specific Ask



Vague ask

Summarize this ticket. Claude has to guess the length, the tone, the fields, and who reads it, so two runs rarely match.




Specific ask

Summarize this ticket in two sentences for an on-call engineer, and list the affected service. Now every run lands in the same shape.

Say What To Do, Not What To Avoid

Positive instructions steer Claude far better than a wall of prohibitions.



Describe the behavior you want
in plain, affirmative terms




Convert each don't into a matching do
the model can follow



Reserve negatives
for genuine hard limits,
not the entire spec

Three Format Levers



Positive framing

Name the behavior you want. Describe the target directly instead of listing everything the model should avoid.



XML tags

Ask for output inside named tags, like prose-paragraph tags, so the shape is unmistakable.



Match your style

Write the prompt in the style you want back. A clean prompt tends to yield clean output.

Add the Why Behind a Rule

Explaining the reason behind an instruction beats a bare command.




Give the motivation, and Claude generalizes to cases you missed



Bare rules only cover the exact case you spelled out



A short why also documents intent for the next developer



Context turns a brittle rule into judgment the model can apply.

Three Ways to Sharpen an Ask




Name the format

Say whether you want JSON, prose, a table, or a list, and name the exact fields the answer should include.



Bound the length

Give a concrete limit, like two sentences or five bullets, so responses stay consistent instead of drifting long.



Sequence the steps

For multi-part work, hand Claude a numbered list of steps so nothing is skipped and the order holds.

When Output Is Inconsistent

Inconsistent output is usually an instruction problem, not a model defect.



Reread the ask

as if you knew nothing about the goal



Tighten the wording

and constraints before touching sampling



Add examples

only after the instruction itself is clear

Clear Instructions, Consistent Output



Make every ask specific, unambiguous, and ordered

Vague prompts create variable outputs. Clarity and order drive consistency.



Tell Claude what to do, not just what to avoid

Positive, actionable guidance yields better and more reliable results.



Add the motivation so the model generalizes your intent

Why it matters helps the model apply the instruction beyond examples.



Name the format, bound the length, sequence the steps

Structure and constraints reduce ambiguity and improve reproducibility.



When output drifts, fix the instruction before sampling


Iterate on the prompt first; sampling tweaks come after the instruction is stable.

System vs User Placement, Constraints

Where you put an instruction matters as much as how you word it.


System-Level Placement

Constraints are stable, global, and hard to override



User-Level Placement

Constraints are easy to override or forget



Three Placement Buckets



System prompt

The place for persistent role, behavior, and constraints. Whatever you set here applies to every turn of the whole session.



User turn

The place for the immediate, per-request input, the actual question, document, or data Claude should act on right now.



Tool result

The place for content a tool returns. Claude is trained to treat it as untrusted data, not as instructions to obey.

System Prompt vs. User Turn



System prompt

Stable across requests. Your role, tone, house rules, and hard constraints, written once so they govern every turn without being repeated.



User turn

Changes every request. The specific input for this one call, the text to summarize or the fields to extract, that would be noise in the system prompt.

Give Claude a Role in System

A role set in the system prompt focuses Claude's behavior and tone.



Even a single sentence
of persona measurably
sharpens output



State the specialty,
the audience, and the
standard to hold



Keep the role stable
so it shapes every response

A clear role up front is one of the cheapest quality gains.



ROLE ASSIGNMENT

TECHNICAL WRITER



Four Placement Rules



Durable rules to system

Standing behavior and constraints go here, written once, not repeated each turn.



Per-request data to user

The specific input for this call belongs in the user message.



Untrusted content to tools

Third-party or fetched content stays in tool results, treated as data.





Your instructions stay with you

Keep your own instructions in a user turn, not a tool result.

Placement Feeds Prompt Caching

Where content sits also decides what prompt caching can reuse.



A stable system prefix can be cached and reused across calls



Volatile user content changes every call, so it can't cache




Durable rules up top keep the cacheable prefix intact

Loose Constraint vs. Precise Constraint



Loose constraint

Respond in JSON. The container is named, but the fields, the nesting, and the extras are left open, so every run decides those details differently.



Precise constraint

Return a JSON object with exactly two fields, summary and service, and nothing outside the object. The shape is fully specified, so nothing is left to guess.

Three Kinds of Output Constraint



Format constraint

Specify the container, JSON, a table, prose, or a list, and the exact fields the answer has to include.



Value constraint

Restrict what an answer may contain, like an allowed set of categories, so responses stay inside a known vocabulary.



Length constraint

Bound the size, two sentences or five bullets, so output stays consistent and cheap to post-process downstream.

Constraints Cut Post-Processing

Every constraint you state up front is code you don't have to write later.



A named format

means no fragile parsing of free text



An allowed value set

means no mapping messy labels back




A length bound

means predictable, uniform responses

CONSTRAINED PROMPT

```
Return JSON matching
this schema:
{
  "title": string,
  "category": one of
    ["billing", "technical",
    "account"],
  "summary": string
    (1 to 120 chars)
}
```

MODEL



CLEAN, STRUCTURED OUTPUT

```
{
  "title": "Payment failed",
  "category": "billing",
  "summary": "User was
  charged twice for the
  same subscription."
}
```



Constraints lock
in structure



Model fills
within bounds



Downstream logic
stays simple

Push the shaping into the prompt, and downstream code stays simple.

Placement and Constraints



Put stable role, behavior, and rules in the system prompt



Keep per-request data in the user turn



Remember placement decides what prompt caching reuses




Constraints shape output but don't guarantee it



Word constraints precisely, and save hard guarantees for structured outputs

Designing Effective Few-Shot Examples

Few-shot examples are one of the strongest levers you have for consistent output.



Zero, One, and Multi-Shot



Zero-shot


Give the task with no worked examples, leaning on pretrained knowledge and a clear instruction. Fits common, well-defined tasks.



...

One-shot

Provide exactly one input-output example. It suits cases where a single canonical example fully defines the pattern and more would add nothing.



Multi-shot

Supply two or more examples to demonstrate format, tone, or reasoning. Diversity and quality matter more than sheer count.

A Deliberate Choice, Not a Ladder

Zero, one, and multi-shot aren't a fallback ladder you climb when the last rung fails.



Match the choice

to task complexity and how specific the format is



A clear zero-shot

can equal few-shot on common tasks



Add examples

only when they improve coverage, not out of habit



Examples that don't earn their tokens are just bloat.

Many Examples vs. Few Good Ones



Many near-duplicates

A big pile of similar examples feels thorough, but it teaches Claude one narrow pattern. The prompt overfits the happy path and misses real variety.



Few diverse ones


A small set of well-chosen examples that span the real input range steers far more reliably. Quality and coverage beat raw count every single time.

Three Marks of a Strong Set



Relevant

Each example mirrors your actual use case closely, so the pattern Claude learns is the one you truly need.



Diverse

Examples vary enough to cover edge cases, so Claude doesn't latch onto an accidental pattern in look-alike samples.




Structured

Each example is clearly separated from your instructions, so Claude reads it as a demonstration, not as a new directive.

The Happy-Path Trap

Examples that only show clean, easy inputs quietly cap how well the prompt handles hard ones.



Cover the cases that break things, not just the easy ones.

Four Ways Example Sets Go Wrong




Too few, too alike

A handful of near-duplicates teaches one narrow pattern and leaves real variety uncovered.



Happy path only

Only clean inputs appear, so the model never learns how the hard cases should look.



Buried in instructions

Examples blur into the directions with no tags, so Claude may read them as commands.



Stale or off-target

Examples drift from the current task or model, quietly steering toward outdated behavior.

Wrap Each Example in Tags

Structure is what lets Claude tell your examples apart from your instructions.



Wrap each example
in its own labeled tags,
and group them together



Three to five examples
hits the sweet spot
for most tasks



Ask Claude to critique
your examples or
draft more

EXAMPLES

`<example_1>`

`</example_1>`

```
{  
  "task": "Summarize this customer review.",  
  "input": "The product arrived quickly and works great.",  
  "output": "Positive review about fast delivery and quality."  
}
```

`<example_2>`

`</example_2>`

```
{  
  "task": "Classify the ticket intent.",  
  "input": "I can't log in to my account.",  
  "output": "intent: account_login"  
}
```

`<example_3>`

`</example_3>`

```
{  
  "task": "Extract key entities from the email.",  
  "input": "Hi, I'm Jane Doe and my order #A123 is late.",  
  "output": {"name": "Jane Doe", "order_id": "A123"}  
}
```

...

(Continue up to 5 examples)

Three Ways to Level Up Examples



Show the reasoning

Put thinking steps inside your examples, and Claude generalizes the reasoning pattern, not just the final answer.



Let Claude help

Ask Claude to evaluate your set for relevance and diversity, or to generate more examples to fill gaps.



Stay canonical

Curate a few diverse, canonical examples that portray the behavior, rather than stuffing in every edge case.

Well-Chosen Examples Win



Choose zero, one, or multi-shot to fit the task, not as a fallback



Keep examples relevant, diverse, and clearly structured with tags




Cover boundary and adversarial inputs, not just the happy path



Let Claude critique and expand your set as you refine it


Refining Prompts Against an Eval Set

A prompt that works once hasn't been validated, it's just been lucky.



EVAL RESULTS (SCORE PER INPUT)			
INPUT	EXPECTED (GOLD)	SYSTEM OUTPUT	SCORE
01			✓ 0.87
02			⚠ 0.42
03			✓ 0.93
04			✗ 0.18
05			✓ 0.76
OVERALL SCORE		NEXT STEP	
 0.63		 Revise prompt to address weak cases.	

The Refinement Loop




Test

Run the prompt across a fixed set of inputs that reflect what production really sends, not one hand-picked example.



Observe

Look at where outputs fail against clear criteria, so you're measuring behavior, not admiring a single good answer.



Adjust

Change the prompt to fix the failures you saw, then run the same set again to confirm the fix.

One Lucky Output Isn't Proof.

Eyeballing a single good output is the most common way prompt work goes wrong.



One passing example
tells you almost nothing
about reliability



Judge the prompt
across representative
inputs, not a favorite case



Include the adversarial
inputs that are built
to break it



Field of Untested Inputs



Short
messages



Long
documents



Other
languages



Code &
snippets



Tables &
structured data



Dates &
time ranges



Locations &
addresses



Emails &
subjects



Product &
sku data



Ambiguous
requests



Edge cases &
boundaries



Repeated &
looped input



Typos &
misspellings



Toxic &
abusive text



Random &
noisy text



Mixed intent
queries



Files &
attachments



Voice to
text



Domain
jargon



Out-of-scope
questions

Three Kinds of Test Input



Representative

The everyday inputs your app really sees in production. They confirm the prompt handles the common case well.



Boundary


The edges, empty fields, huge inputs, ambiguous requests. They expose where behavior starts to bend.



Adversarial


Inputs crafted to trip the prompt or smuggle in instructions. They test whether it holds under pressure.

Eyeballing vs. a Fixed Eval Set



Eyeballing outputs

You tweak the prompt, glance at one result, and call it good. Each change is judged on a different input, so you can't tell what actually improved.



A fixed eval set

You run every version against the same frozen inputs. Now a change is measurable, comparable, and safe to keep or roll back.

Self-Correction as a Chain

One powerful refinement pattern runs Claude over its own work in stages.

Draft Review Refine Chain



API Call 1 Draft Answer

model: claude-opus-5



API Call 2 Review vs. Criteria

model: claude-opus-5



API Call 3 Refine Answer

model: claude-opus-5



Draft the answer in a first call

Claude produces an initial draft to establish a baseline response.



Review it against explicit criteria in a second call

Claude evaluates the draft using clear criteria and returns a targeted review with strengths and gaps.



Refine based on that review in a third

Claude applies the review to improve accuracy, clarity, and completeness in a final refined answer.

A Cheap Self-Check

Sometimes the whole fix is one extra line at the end of the prompt.




Ask Claude to verify its answer against the criteria before finishing



This catches errors reliably, especially in code and math



It costs one sentence and adds a built-in second look



A self-check turns a single pass into a quick draft-and-review.

Prompts Are Versioned Artifacts



Version it

Track prompts in source control like code, so every change is reviewable, reproducible, and safe to roll back.



Pin the model

Tie the prompt to a specific model ID, since a new model release can shift behavior your prompt relied on.




Regression-test it

Re-run the prompt against your frozen eval set on every change, so a fix doesn't quietly break another case.

Validated Means Tested

A change that passed on one happy-path example is not a validated change.



Treat every prompt edit like a code change that can regress

Even small edits can break assumptions or edge-case behavior.



Gate it behind the same eval set you test everything on

Use the full set, including hard and adversarial cases.



Ship only when the whole set, hard cases included, passes

One failure means the change is not validated.

Validation is a bar the prompt clears, not a single lucky run.

Refine Against Real Inputs



Refine in a loop, test, observe failures, then adjust



Measure against a fixed set of representative and adversarial inputs





Use self-correction chains and a self-check step to catch errors



Version, pin the model, and regression-test before shipping

Sanitizing Untrusted Input

Untrusted input is any content your application did not author.
In the prompt it can smuggle instructions the model may follow.



Three Core Terms



Input Sanitization

Cleaning, validating, or delimiting user-supplied content before it enters the prompt, so the model treats it as data, not instructions.



Untrusted Input

Any content the application did not author: user messages, retrieved documents, and tool results. Treat all of it as adversarial.



Content Boundaries

Explicit delineation of trusted developer instructions from untrusted content, often by wrapping data in labeled tags.

What Sanitization Does

Sanitization is prep work you run before the model call.



Clean the input, stripping or escaping characters that could break your structure



Validate it against expected shape, type, and length



Delimit it, wrapping the content so its edges are unambiguous

INPUT PIPELINE: RAW - READY



Done well, the model reads the payload as data, never as a fresh instruction.

Data Versus Instructions



Trusted Instructions

Your system prompt and developer instructions define what the model should do. They are authored by you and carry authority over the task.



Untrusted Data

External text, user fields, and tool results the model works on. It should never be read as a command, only as material to process.

Wrapping Input in Labeled Tags

The most reliable delimiter is a labeled tag around the untrusted block.



Instructions (outside tag)

Summarize the report for Emily.
Focus on risks and actions.



Opening tag

`<user_document>`



Untrusted content (inside tag)

```
Q2 Security Review - Acme Corp\nDate: 2025-04-15\nFindings:\n- Outdated dependencies in auth service\n- Excessive IAM permissions\n- Missing rate limits on signup API\nActions:\n- Upgrade libraries\n- Apply least privilege\n- Add rate limiting middleware\nOwner: Platform Team
```

`</user_document>`



Closing tag



Instructions (outside tag)

If the report implies urgent risk, say so.
End with next steps.



Open and close

the content with a
clear tag like `user_document`



Tell the model


that anything inside
those tags is data



Keep your own

instructions outside
the tags, never mixed in

Where Untrusted Content Belongs




Tool Results

Deliver third-party content inside tool_result blocks, where the model is trained to treat instructions with skepticism.



Not the System Prompt

Never paste untrusted text into the system prompt or a plain instruction, where it inherits authority.



Your Instructions Separate

Send your own directions in your turn, not buried in the same block as untrusted data.

Validating Structured Fields

Not all untrusted input is free text. Structured fields get their own check.



Validate each field's
type, range, and allowed
values before it reaches
the prompt



Reject or escape
anything that fails
instead of passing
it through



**Run this on the
developer side,**
in code, before
the model call



Where It Sits in the Path



Before the Call

Sanitization runs on the developer's side, in your code, before the request ever reaches the model.



Alongside Schema Checks

It works next to schema validation of structured fields, one handling free text, the other typed values.




A Reduction, Not a Cure

Delimiting and validating lower the risk of smuggled instructions but do not fully stop a determined injection.

Escaping to Prevent Breakout

A determined input tries to break out of your wrapper and start a new instruction.



Clean edges keep smuggled text from posing as your command.

Treating Data as Data



Clean, validate, and delimit content before the prompt



Wrap untrusted blocks in labeled tags, instructions outside



Keep untrusted content in tool results, not the system prompt



Validate structured fields by type, range, and allowed values



Remember it reduces injection risk, it doesn't fully cure it

Getting Reliable Structured Output

Not every way of asking for JSON is equally reliable.

Shapes the Output

Best effort, never a guarantee

1



No Constraints

Let the model decide.
Highest risk of invalid output.

2



Format Instruction

Ask in prose for JSON only.
Often works, still best effort.

3



Few-Shot Examples

Show sample input-output pairs.
Steers strongly, guarantees nothing.

Enforces the Schema

Constrained decoding, guaranteed

4



Define the Schema

Fields, types, required flags, enums.
The contract output must match.

5



Structured Outputs

Attach the schema to the request.
Whole response matches the schema.

6



Strict Tool Use


Turn on strict mode for the tool.
Tool input is grammar-enforced.

Three Rungs on the Ladder




Format instructions

A prose request like return only JSON. It shapes the output and often works, but it stays best effort, never a guarantee.



Few-shot examples

Sample input-output pairs that demonstrate the target shape. They steer output strongly, yet on their own still guarantee nothing.



Schema enforcement

Structured outputs or strict tool use. Constrained decoding guarantees the response matches the schema you defined.

Shape Is Not the Same as Guarantee

Syntactically valid JSON and schema compliance are two different bars.

A response can parse cleanly yet miss a required field

JSON may be well-formed, but if a required field is absent, your code can still break.

It can carry the right keys with the wrong types

Types matter. A string where a number is expected can pass JSON but fail validation.

It can invent an enum value your code never expected

Enums exist to protect your system from unexpected values that lead to bugs.

So pick the mechanism the requirement demands, not the one that's easiest.

Prose Request vs. Enforcement



Prose request

You ask for JSON in the instruction. The model usually complies, but nothing stops a stray sentence or a missing field, so your parser can still break.




Enforcement

You attach a schema via structured outputs or strict tool use. Constrained decoding makes the shape mandatory, so valid, typed output is guaranteed.

Two Ways to Enforce a Schema

When you need a hard guarantee, you have two mechanisms to reach for.




JSON outputs, the whole-response form of structured outputs, use a `json_schema` via `output_config.format`



Strict tool use sets `strict` to `true` so a tool's input schema is grammar-enforced

JSON Outputs vs Strict Tool Use



JSON Outputs

Set `output_config.format` to a `json_schema` and the entire response is constrained to that shape.

Use it when the model's final answer itself has to be a structured object.




Strict Tool Use

Set `strict` true on a tool and the tool's `input_schema` is guaranteed.

Use it when Claude calls a function and the arguments have to be valid.

Schema Design Essentials




Precise names and types

Name each field clearly and pin its type. Vague names invite wrong or ambiguous values.



Required vs optional

Mark truly needed fields required. Note that `additionalProperties` has to be false on objects.



Enums and descriptions

Constrain categorical fields with an enum, and describe each field so values stay unambiguous.

What Constrained Decoding Buys You

Enforcing the schema removes a whole class of runtime failures.




Output is always valid JSON,
so no parse errors to catch



Fields come back type-safe,
with required keys present



No retry loop
just to fix a schema violation



You spend your validation budget on meaning, not on syntax.

When to Climb the Ladder




Fields must hold

If downstream code depends on specific fields and types, reach for structured outputs, not a prose ask.



Enums must hold

When a value has to come from a fixed set, strict tool use or an enum in the schema locks it in.



Convenience is not the test

Pick the most reliable mechanism the requirement demands, even when a prose request would be quicker to write.

Prefill Is No Longer the Answer

Older code forced JSON by prefilling the start of the reply. That path is closed.

Deprecated: Prefill JSON Start



User

Summarize this incident.



Assistant (prefill)

```
{ "summary": "  
  "severity": "  
  ...
```



**400
Bad Request**

Supported: Structured Outputs



User

Summarize this incident.



Schema

```
summary: string  
severity: string  
impacted_services: array
```



Assistant (JSON)

```
{  
  "summary": "Database outage resolved",  
  "severity": "High",  
  "impacted_services": ["DB", "API"]  
}
```



On Claude 4.6 and later,
a prefilled final turn
returns a 400 error



**The supported replacement
is structured outputs,**
not a workaround



Enforce the shape
with a schema instead of
steering the first tokens

Climbing the Reliability Ladder



Instructions and examples shape output but guarantee nothing



Structured outputs and strict tool use enforce the schema



Valid JSON is not schema compliance




Design precise fields, types, required flags, and enums



On 4.6 and later, enforce format with schemas

Validate and Re-Prompt on Failure

Even guaranteed structure can be plausibly wrong.



Three Validation Ideas



Response validation

Programmatically checking a reply against a schema and your business rules after generation, covering meaning, not just syntax.



Semantic check

Confirming the values make sense and are complete, not merely that the JSON parsed without error.



The correction path

On failure, re-prompting with the specific error so the model can fix its own output.

The Validate and Re-Prompt Loop

Reliable pipelines treat generation as one step in a loop, not a finish line.





Generate the response
from the model



Validate it
against the schema and
your business rules



On failure, re-prompt
with the exact error so
the model corrects



Validation Failed



A one-shot check with no correction path is where reliability leaks out.

Syntax Check vs. Semantic Check



Syntax check

Did the response parse and match the schema shape? Necessary, sure, but it says nothing about whether the actual values are correct.




Semantic check

Do the values obey your business rules and make sense together? This is the check that catches a plausible but wrong answer before you act.

Why One-Shot Checks Fail

A single check at the end assumes the first generation is either right or throws an error.



Silent failures

return success with wrong or incomplete data



No exception fires,

so bad data flows downstream untouched



Non-determinism

means the same prompt can vary run to run

Four Structured-Output Edge Cases



Refusal

If the model refuses on safety grounds, `stop_reason` is `refusal` and the output may not match the schema.




Max tokens

If generation hits the token limit, `stop_reason` is `max_tokens` and the object can be truncated.



Enum capitalization

Enum string values can differ only in case, so compare them case-insensitively, not exactly.



Property ordering

Required properties come first, then optional, regardless of your schema's declared order.

Branch on stop_reason First

Before you trust a structured payload, check why the model stopped.




A refusal means no schema-valid object, so handle it as a refusal



A max_tokens stop means the data may be cut off midway



Only an end_turn stop earns your normal parsing path



Reading stop_reason first keeps you from parsing an answer that was never finished.

Three Validation Layers



Schema validation

Confirm the response parses and every field has the declared type. This is the syntax gate.



Business-rule checks

Confirm the values are allowed, consistent, and complete for your domain, beyond what the schema can express.



Re-prompt on failure

Feed the exact validation error back so the model corrects itself, closing the loop.

Design for the Re-Prompt

A useful re-prompt gives the model everything it needs to fix the specific problem.



Re-Prompt Example

Carry the exact error back to the model.



1) Context

Goal: Get current weather for a location
Tool: get_weather
Parameters: { "location": "San Francisco" }



2) Exact Error

Error Type: ValidationError
Message: 'location' is required
When: 2025-05-12T14:32:18Z



3) Offending Output

```
{  
  "type": "tool_use",  
  "name": "get_weather",  
  "input": {}  
}
```



4) Retry Policy

Retries So Far: 2
Max Retries: 3
Action if Exceeded: Escalate to human review



5) Instructions

Provide a corrected tool call that includes the required 'location' parameter and nothing else. Do not repeat the same error.



6) Response Format

Return only a JSON object with keys: type, name, input. No explanations.



Include the exact error, not just a generic try again

Give the model the precise error type, message, and when it occurred so it can reason about the failure correctly.



Show the offending output so the model sees what it produced

Include the exact tool call, response, or text that caused the error. Don't make it guess.



Cap the retries so a stubborn failure escalates instead of looping

State the retry count, the maximum, and what should happen next if the limit is reached.

Distinguish a genuine refusal from a bug, since each needs a different fix.

Generate, Validate, Re-Prompt



Validate meaning and completeness, not just JSON syntax



Re-prompt with the exact error so the model can correct



Structured output still fails on refusal and max_tokens



Check stop_reason and compare enums case-insensitively



Non-determinism means every run needs the same validation

Parsing Model Output Defensively

Model output looks like a clean contract, right up until a generation comes back malformed, truncated, or missing a field.



Model

Raw Model Output

```
{
  "name": "Acme Road Runner",
  "status": "active",
  "plan": "Pro",
  "items": [
    { "id": "tool_1", "qty": 2 },
    { "id": "tool_2", "qty": "two" },
  ],
  "total": "19.99",
  "discount": 5,
  "notes": "Thanks!",
  "unexpected": { "foo": "bar" }
  "incomplete": true
  ...
}
```



**Malformed,
truncated, or
missing fields**



Defensive Parsing Pipeline



JSON Parse

fail



Reject Output

pass



Schema
Validation

fail



Reject Output

pass



Type Check

fail



Reject Output

pass



Field Completeness

fail



Reject Output

pass



Sanitize & Coerce

pass



Typed Object

Three Habits of Defensive Parsing



Tolerant extraction

Pull the fields you need without assuming the whole payload is well-formed. Read what's there, and cope when something is missing.



Explicit failure handling

Decide up front what happens on bad output. A wrapped parse, a fallback, or a re-prompt, never an unhandled crash.



No blind trust

Treat every field as possibly absent or mistyped, so one bad generation degrades gracefully instead of taking the pipeline down.

Never Assume a Field Is There

A response that parsed as JSON can still lack the shape your code expects.



Confirm each required field exists before you read it

Presence checks prevent crashes and give better error messages.



Check the type before doing math or string work

Type mismatches cause subtle bugs and runtime exceptions.



Treat an empty or null value as an expected case

Gracefully handle missing or empty data in your logic.

Verify each field before using the data



All checks pass → Safe to use data

Trusting Output vs. Parsing It



Trust it

You read fields straight off the response and act on them. One malformed or missing value throws deep in your logic, and the whole flow crashes on a single bad generation.



Parse it

You extract through checks that expect surprises. A bad value is caught at the boundary, logged, and handled, so one odd response degrades instead of taking everything down.

Drill: Extractor That Crashed Overnight

Let's apply the last lesson's edge-case roster to a real debugging scenario.

Ticket Triage Extractor Pipeline

02:17 AM



extract.py:87 extract_fields()

parser.py:142 parse_ticket_json()

helpers.py:204 get_priority()

KeyError: 'priority'

Raw LLM Response (truncated)

```
1 {
2   "ticket_id": "TCK-98231",
3   "summary": "Login page error",
4   "category": "Authenti
```



Response ended abruptly



A ticket-triage extractor runs clean for days, then crashes overnight



The error is a **KeyError** on priority, three functions deep



The raw response ends mid-object, with the closing braces missing



The crash site isn't the cause. The clues point back to the response itself.

Walking the Diagnosis



Read stop_reason first

The response reports max_tokens, so generation hit the ceiling and the object was cut off mid-structure.



Inspect the raw payload

A long ticket blew past the token budget, leaving truncated JSON that no parser could finish.




Name the missing guard

The code parsed before checking why the model stopped, so a known edge case landed as a mystery crash.

The Hardened Version

The fix moves every check to the boundary, before any field is read.



Branch on stop_reason first,
and route refusal and cutoff
to their own paths



Raise the token budget
and guard each field
access by name



Compare enum values
case-insensitively, per the
roster from last lesson

Guarded Parsing Boundary

Catches a truncated response cleanly

Raw Response

```
{ "user": "u123",  
  "items": [ { "id": "a1",  
               "name": "Widget", "qty": 2 },  
             { "id": "b2", "name": "Gadget",  
               "qty": 1 }    <-- cutoff
```

Check stop_reason



Refusal Path



Cutoff Path

Field Guards (by name)

user



Exists?

items



Exists?

id



Exists?

qty



Exists?

Parsed Output

Valid, complete data




Three Safe Responses to Bad Output



Fail loud

When output can't be parsed, raise a clear error at the boundary instead of passing garbage downstream to fail somewhere confusing.



Fall back

Have a default path, a safe value, a cached result, or a degraded mode, so one bad generation doesn't stall the user.



Re-prompt

Send the response back with the error described, giving the model a chance to correct what it got wrong.

One Failure Parsing Can't Catch

Every guard so far catches output that is visibly broken.

</> Service Response (200 OK)

```
1 {
2   "status": "success",
3   "answer": "The capital of France is Berlin.",
4   "country": "France",
5   "confidence": 0.97,
6   "metadata": {
7     "model": "claude-opus-5",
8     "latency_ms": 642
9   }
}
```



What's really wrong

The answer is actually: **Paris**.

The response looks perfect to any parser or schema validator.



API Response



Schema Valid



All Guards Pass



Wrong Answer Shipped



A silent failure

returns success while the content is wrong



It parses cleanly,

so no guard in this lesson will flag it



Catching it takes

a semantic gate, judged on meaning

Parse Like You Expect Surprises



Extract tolerantly and never assume a field is present



Check the stop reason first, since refusal or cutoff breaks the shape



Guard structured output too




Choose a safe response, fail loud, fall back, or re-prompt



Silent failures need a semantic gate, covered next lesson

Distrusting Confident-But-Wrong Output

A model states a wrong answer with the exact same fluency as a right one.




Three Truths About Confidence



Confidence isn't correctness

A fluent, self-assured answer is not evidence it's right. Tone tells you nothing about accuracy.



Wrong sounds like right

Models state wrong answers with the same certainty as correct ones. There's no built-in hedge on a mistake.



Verify from outside

Correctness comes from grounding, citation, or validation, never from how sure the response sounds.

Why Fluency Fools You

Fluent prose reads as authority, but generation doesn't work that way.




The model predicts plausible text, not verified fact

It optimizes the next token based on patterns, not on whether a claim is true in the real world.



A confident tone is a style, decoupled from ground truth

Polished, certain language comes from training, not from access to external verification.



Detail and specificity can make a wrong answer look more credible

Invented specifics, names, and numbers mimic correctness and often go uncorrected.



Train yourself to separate how an answer sounds from whether it's actually true.

Sounds Right vs. Is Right



Sounds right

The answer is fluent, detailed, and self-assured. That style is generated the same way whether the content is true or false, so it proves nothing on its own.



Is right

The claim is backed by a source you can check, a computed result, or a validation gate. Correctness is established outside the model, not inferred from tone.

Ground Before You Trust

For anything that matters, make the model show its work against something real.



Claim to Verify

Tokyo is the capital of Japan.



Model Reasoning

I will verify this using a trusted source.



Tool Call

Get Wikipedia page summary for Tokyo



Observation

Source: wikipedia.org - Fetched: 2025-05-24



Retrieved Evidence

"Tokyo is the capital and most populous city of Japan."



Verified Answer

Tokyo is the capital of Japan.



Have it cite the source or document it drew the claim from

Require the model to show where the information came from, with names, links, or identifiers.



Prefer a computed or looked-up value over a recalled one

Favor fresh tools, APIs, and databases instead of relying on the model's memory or training data.



Ask it to self-check the answer against your criteria before finishing


Make the model evaluate its own answer against your rubric and fix any issues.

Four High-Stakes Guardrails




Ground the claim

Require a citation or a source Claude can point to, so the answer rests on something checkable.



Cross-check

Verify a critical value a second way, a computation, a lookup, or a second pass, before acting.



Validate semantics

Run a gate that judges whether the output is complete and plausible, not just well-formed.



Gate the action

Put a human or a check between confident output and any irreversible downstream action.

The Silent Failure Trap

The dangerous failures are the ones that return success while being wrong.




No exception fires,
so try/except never
catches it



The status code is fine,
so error handling
waves it through



Only a semantic check
on the content itself
can flag it



Treat a clean status as necessary but never sufficient for a correct result.

Three Verification Sources



Grounding

Tie the answer to real context the model was given, so a claim traces back to provided data, not memory.



Citation

Require the response to name where a fact came from, making it checkable instead of asserted.




Validation

Programmatically test the output against a schema or business rule, catching wrong-but-fluent answers.

Close the Loop on High-Stakes Output

Skepticism only pays off if it's wired into the flow, not left as a good intention.



Decide which outputs are high-stakes
and demand grounding there



Insert the validation gate
before the output drives any action



Route a failed check
to a fallback or a human,
not to silent acceptance

Verify, Don't Trust the Tone



A wrong answer sounds exactly as sure as a right one



Confidence does not reliably signal accuracy in generation



Ground claims with sources, citations, or computed values



Add a semantic gate to catch silent, wrong-but-successful output



Verify before confident output drives any downstream action