



Complete Video Tracker — 159 Videos

Playlist: Abhishek Veeramalla — DevOps Engineer in 3 Months (2026) [YouTube Playlist](#) **Total**
Videos: 159 | **Schedule:** Jul 31 – Dec 31, 2026 (≈1.2 videos/day)



Overall Progress Dashboard

Phase	Videos	Est. Hours	Date Range	Status
1: Foundations	36	~31 hrs	Jul 31 – Sep 4	
2: Containers & K8s	45	~37 hrs	Sep 5 – Oct 19	
3: CI/CD & IaC	35	~27 hrs	Oct 20 – Nov 23	
4: Advanced Topics	30	~23 hrs	Nov 24 – Dec 23	
5: AI & Career	13	~10 hrs	Dec 24 – Dec 31	
TOTAL	159	~128 hrs	Jul 31 – Dec 31	

Legend: Not Started | In Progress | Complete



How to Use This Tracker

- ☐ → ☒ Check off each video as you watch it
- Add brief notes in the Notes column for key takeaways
- Review progress every Sunday evening
- If you fall behind, skip interview Q&A videos and revisit later
- If you're ahead, do extra hands-on labs

Note: Video titles and durations are approximate based on the playlist curriculum structure. The actual playlist is dynamic — verify titles as you watch. Durations are estimated at ~40–65 min per video based on course averages.

2027 HORIZON: AI PLATFORM ENGINEERING (MLOps & LLMOps)

(Once you conquer the core Enterprise DevOps curriculum by Dec 31st, 2026, this is your precise roadmap to evolve into a highly specialized AI Platform Engineer.)

Phase 1: Generative AI & Developer Productivity

- Advanced Prompt Engineering:** Precision generation of Terraform configurations, bash scripting, and regex via LLMs.
- AI Coding Assistants:** Mastering GitHub Copilot CLI, Cursor, and Codeium for rapid pipeline architecture.
- AI-Assisted CI/CD:** Using LLMs to auto-diagnose Kubernetes crash loops and GitLab pipeline failures.

Phase 2: AIOps (Artificial Intelligence for IT Operations)

- **Predictive Observability:** Datadog Watchdog & Dynatrace Davis AI setup for anomaly detection prior to system impact.
- **Automated RCA:** Feeding ELK/Splunk logs into anomaly engines for zero-touch Root Cause Analysis.
- **AI-Driven FinOps:** Dynamic cloud resource autoscaling and right-sizing via Kubecost AI and AWS Compute Optimizer.

Phase 3: MLOps (Machine Learning Operations)

- **Data Versioning:** Implementing DVC (Data Version Control) over AWS S3.
- **ML CI/CD Pipelines:** Building model delivery pipelines via Kubeflow, MLflow, or AWS SageMaker block definitions.
- **Model Serving APIs:** Deploying trained inference models to production utilizing KServe and TensorFlow Serving on EKS.

Phase 4: LLMOps & GPU Platform Engineering

- **GPU Orchestration:** The pinnacle skill—provisioning and managing NVIDIA GPU node groups in Kubernetes natively.
- **Vector Database IaC:** Deploying and scaling Milvus, Pinecone, or pgvector infrastructure.
- **Open-Source LLM Hosting:** Containerizing local massive models (Llama 3, Mistral) via vLLM and HuggingFace TGI on Kubernetes clusters.

Phase 5: DevSecOps for AI

- **OWASP for LLMs:** Mitigating Prompt Injections, Model Denial of Service (DoS), and Data Poisoning via Web Application Firewalls.
- **Data Privacy Guardrails:** Air-gapping VPCs and configuring Azure OpenAI Private Links to prevent corporate IP leakage.